

Image Processing


INT3404 20

Lecturer: Nguyen Thi Ngoc Diep, Ph.D.

Email: ngocdiep@vnu.edu.vn

Slide & code: https://github.com/chupibk/INT3404_20

Recall week 3: Histogram



An image with L-level intensities

r_k : intensity level k ($k = 0, 1, 2, \dots, L-1$)

n_k : number of pixels with intensity r_k

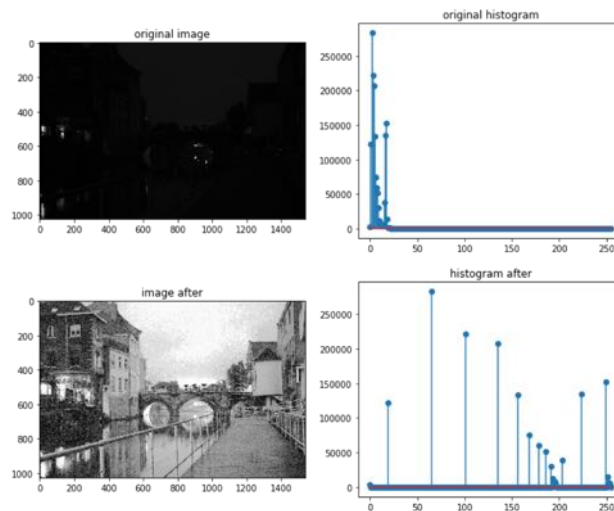
Unnormalized histogram:

$$h(r_k) = n_k$$

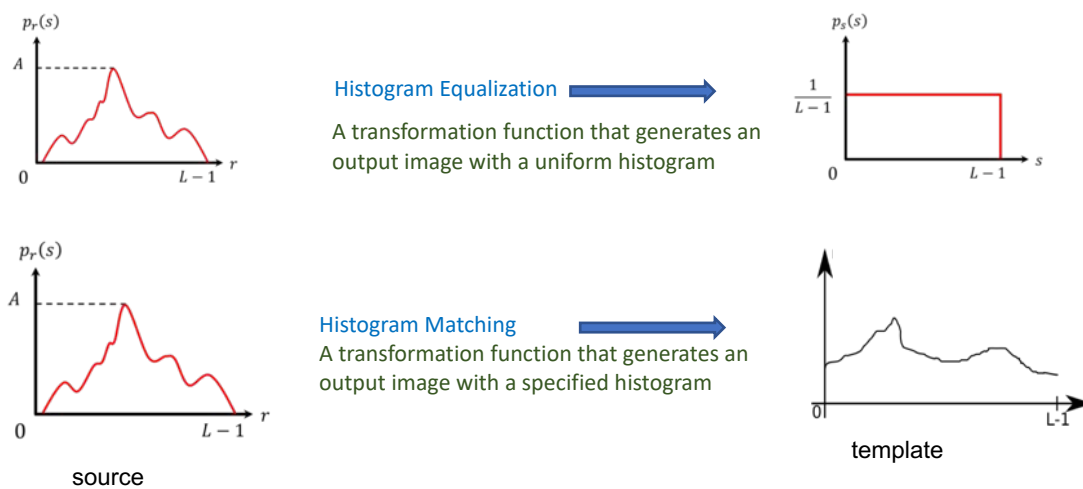
Normalized histogram:

$$p(r_k) = \frac{h(r_k)}{MN} = \frac{n_k}{MN}$$

Recall week 3: Histogram and image appearance



Recall week 3: Histogram equalization and matching




Schedule

Week	Content	Homework
1	Introduction	Set up environments: Python 3, OpenCV 3, Numpy, Jupyter Notebook
2	Digital image – Point operations Contrast adjust – Combining images	HW1: adjust gamma to find the best contrast
3	Histogram - Histogram equalization – Histogram-based image classification	Self-study
4	Spatial filtering - Template matching	Self-study
5	Feature extraction Edge, Line, and Texture	Self-study
6	Morphological operations	HW2: Barcode detection → Require submission as mid-term test
7	Filtering in the Frequency domain Announcement of Final project topics	Final project registration
8	Color image processing	HW3: Conversion between color spaces, color image segmentation
9	Geometric transformations	Self-study
10	Noise and restoration	Self-study
11	Compression	Self-study
12	Final project presentation	Self-study
13	Final project presentation Class summarization	Self-study


5

Week 4: Spatial filtering

Neighbors of a pixel



4 - neighbors



8 - neighbors

Distance between two pixels (1/2)

2 pixels $p=(x, y)$ and $q=(u, v)$

Euclidean distance: $D_e(p, q) = \left[(x-u)^2 + (y-v)^2 \right]^{\frac{1}{2}}$

City-block distance: $D_4(p, q) = |x-u| + |y-v|$

All pixels that are less than or equal to some value d form a diamond centered at (x, y)

Example:

$$\begin{array}{ccc} & 1 & \\ 1 & 0 & 1 \\ & 1 & \end{array}$$
 $D_4 = 1 \ (\rightarrow 4 \text{ neighbors})$

$$\begin{array}{ccccc} & & 2 & & \\ & 2 & 1 & 2 & \\ 2 & 1 & 0 & 1 & 2 \\ & 2 & 1 & 2 & \\ & & 2 & & \end{array}$$
 $D_4 = 2$

Distance between two pixels (2/2)

2 pixels $p=(x, y)$ and $q=(u, v)$

Chessboard distance: $D_8(p, q) = \max(|x - u|, |y - v|)$

All pixels that are less than or equal some value d form a square centered at (x, y)

Example:

```
1 1 1
1 0 1
1 1 1
```

$D_8 = 1$ (\rightarrow 8 neighbors)
square size: 3x3

```
2 2 2 2 2
2 1 1 1 2
2 1 0 1 2
2 1 1 1 2
2 2 2 2 2
```

$D_8 = 2$
square size: 5x5

Spatial filter kernel

- Also called: mask, template, window, filter, kernel
- A kernel: an array whose **size defines the neighborhood of operation**, and whose **coefficients determine the nature** of the filter
- Spatial filtering modifies an image by replacing the value of each pixel by a function of the values of the pixel and its neighbors


Linear spatial filtering mechanism

- A linear spatial filter performs a sum-of-products operation between an image f and a filter kernel, w
- Kernel center $w(0,0)$ aligns with the pixel at location (x,y)

Kernel size: $m \times n$
 $m = 2a + 1$
 $n = 2b + 1$
 Image size: $M \times N$

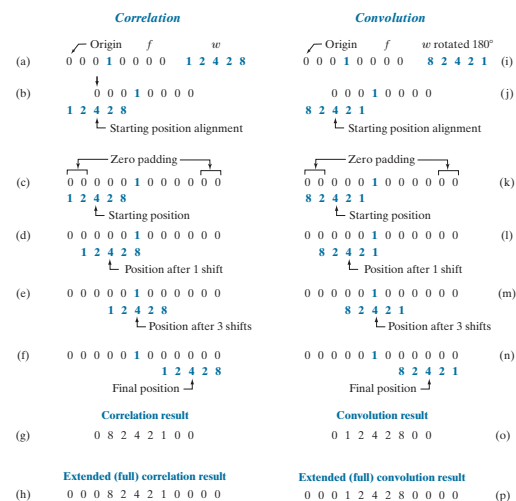
Linear spatial filtering:

$$g(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)$$



Source: Fig. 3.28, Gonzalez

Spatial correlation and convolution in 1D



At the location of the impulse (1)

Source: Fig. 3.29, Gonzalez

Spatial correlation and convolution in 1D

	<i>Correlation</i>	<i>Convolution</i>
	<p>Origin f w</p> <p>0 0 0 1 0 0 0 0 1 2 4 2 8</p>	<p>Origin f w rotated 180°</p> <p>0 0 0 1 0 0 0 0 8 2 4 2 1</p>
At the location of the impulse (1)	<p>Correlation result</p> <p>0 8 2 4 2 1 0 0</p> <p>yield a copy of w, but rotated by 180°</p>	<p>Convolution result</p> <p>0 1 2 4 2 8 0 0</p> <p>yield a copy of w, by pre-rotating w before performing shifting/sum-of-products</p>

Source: Fig. 3.29, Gonzalez

Correlation and Convolution in 2D

	Padded f	
Origin f	0 0	
w	0 0 0 0 1 0	
(a)	(b)	
Initial position for w	Correlation result	Full correlation result
1 2 3 4 5 6 7 8 9 0 0 0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0 0 9 8 7 0 0 6 5 4 0 0 3 2 1 0 0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 9 8 7 0 0 0 0 6 5 4 0 0 0 0 3 2 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
(c)	(d)	(e)
Rotated w	Convolution result	Full convolution result
9 8 7 6 5 4 3 2 1 0 0 0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0 0 1 2 3 0 0 4 5 6 0 0 7 8 9 0 0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 2 3 0 0 0 0 4 5 6 0 0 0 0 7 8 9 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
(f)	(g)	(h)

Source: Fig. 3.30, Gonzalez

Correlation vs Convolution

Correlation


$$(w \star f)(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)$$

Convolution

$$(w \star f)(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x - s, y - t)$$

Correlation vs convolution

$$(w \star f)(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)$$



$$(w \star f)(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x - s, y - t)$$

Fundamental properties of convolution and correlation

Property	Convolution	Correlation
Commutative	$f \star g = g \star f$	—
Associative	$f \star (g \star h) = (f \star g) \star h$	—
Distributive	$f \star (g + h) = (f \star g) + (f \star h)$	$f \star (g + h) = (f \star g) + (f \star h)$

Boundary issues




Image size: $M \times N$

Kernel size: $m \times n$

Output:

$(M + m - 1) \times (N + n - 1)$

$M \times N$

$(M - m + 1) \times (N - n + 1)$

What to do around the borders


- Pad a constant value (black)
- Wrap around (circulate the image)
- Copy edge (replicate the edges' pixels)
- Reflect across edges (symmetric)




Template matching

Template matching


- What if we cut little pictures out from an image, then tried to do cross correlation them with the same or other images?



Template matching



Template matching



Actually...

I subtracted the mean gray value from both the image and the template before doing cross correlation.

Why?

Problem with correlation of raw image template

Consider correlation of template with an image of constant grey value:

a	b	c
d	e	f
g	h	i

 \otimes

v	v	v
v	v	v
v	v	v

Result: $v \cdot (a+b+c+d+e+f+g+h+i)$

Problem with correlation of raw image template

Now consider correlation with a constant image that is twice as bright

a	b	c
d	e	f
g	h	i

 \otimes

2v	2v	2v
2v	2v	2v
2v	2v	2v

Result: $2 \cdot v \cdot (a+b+c+d+e+f+g+h+i)$
 $> v \cdot (a+b+c+d+e+f+g+h+i)$

Larger score, regardless of what the template is!

Solution

- Subtract off the mean value of the template
- In this way, the correlation score is higher only when darker parts of the template overlap darker parts of the image, and the brighter parts of the template overlap brighter parts of the image

Correspondence problem

Finding corresponding feature across two or more views

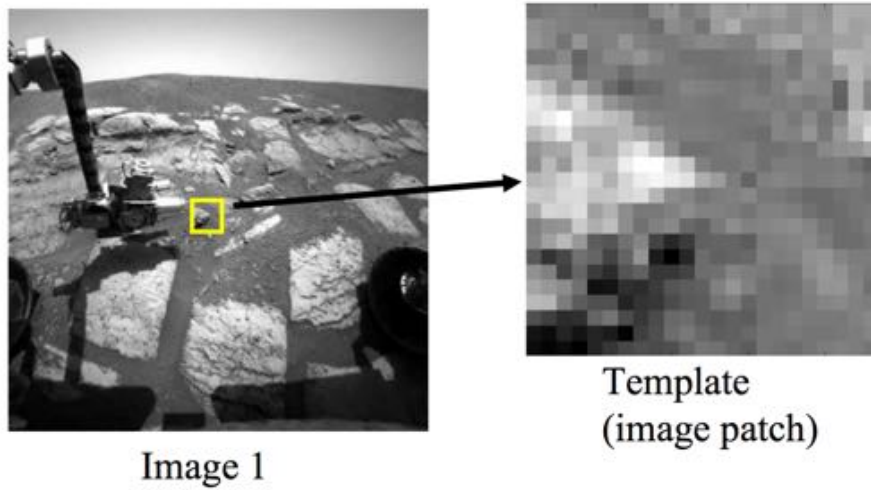


Image 1

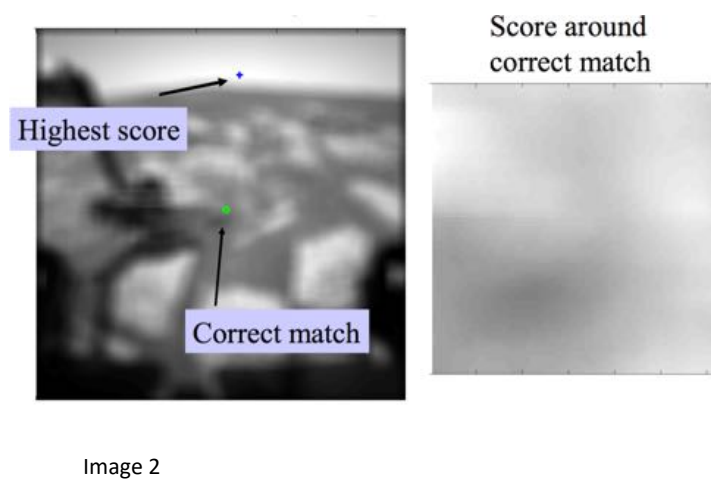
Image 2

Note: this is a stereo pair from the NASA Mars rover.
The rover is exploring the "El Capitan" formation


Example



Example: Raw cross-correlation



Example: Cross-correlation with zero-mean template



Better! But highest score is still not the correct match.
Note: highest score IS best within local neighborhood of correct match.

“SSD” or “Block matching” (Sum of squared differences)

$$\sum_{[i,j] \in R} (f(i,j) - g(i,j))^2$$

- 1 – The most popular matching score
- 2 – T&V claim it works better than cross-correlation


Relation between SSD and Correlation

$$\begin{aligned} SSD &= \sum_{[i,j] \in R} (f - g)^2 \\ &= \sum_{[i,j] \in R} f^2 + \sum_{[i,j] \in R} g^2 - 2 \left(\sum_{[i,j] \in R} fg \right) \end{aligned}$$

$$C_{fg} = \sum_{[i,j] \in R} f(i,j)g(i,j)$$

Correlation!

SSD



Best match (highest score) in image coincides with correct match in this case!

Handling intensity changes

- the camera taking the second image might have different intensity response characteristics than the camera taking the first image
- Illumination in the scene could change
- The camera might have auto-gain control set, so that it's response changes as it moves through the scene.

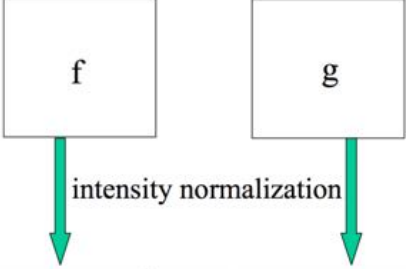


Intensity normalization

- When a scene is imaged by different sensors, or under different illumination intensities, both the SSD and the $C_{\{fg\}}$ can be large for windows representing the same area in the scene!
- A solution is to NORMALIZE the pixels in the windows before comparing them by subtracting the mean of the patch intensities and dividing by the std.dev.

$$\hat{f} = \frac{f - \bar{f}}{\sqrt{\sum (f - \bar{f})^2}} \quad \hat{g} = \frac{g - \bar{g}}{\sqrt{\sum (g - \bar{g})^2}}$$


Normalized cross correlation



$$\hat{f} = \frac{f - \bar{f}}{\sqrt{\sum (f - \bar{f})^2}} \quad \hat{g} = \frac{g - \bar{g}}{\sqrt{\sum (g - \bar{g})^2}}$$

$$\text{NCC}(f,g) = C_{fg}(\hat{f}, \hat{g}) = \sum_{[i,j] \in R} \hat{f}(i,j) \hat{g}(i,j)$$

Normalized cross correlation



Highest score also coincides with correct match.
Also, looks like less chances of getting a wrong match.

Normalized cross correlation

- Important point about NCC:
 - Score values range from 1 (perfect match) to -1 (completely anti-correlated)
- Intuition: treating the normalized patches as vectors, we see they are unit vectors. Therefore, correlation becomes dot product of unit vectors, and thus must range between -1 and 1

Spatial filter kernels

Filter design

- Based on mathematical properties
 - E.g.: A filter that computes the average of pixels in a neighborhood blurs an image
 - E.g.: A filter that computes the local derivative of an image sharpens the image
- Based on sampling a 2D spatial function whose shape has a desired property
 - E.g.: Samples from a Gaussian function to construct a weighted-average filter
- Based on a specified frequency response

Smoothing filters



- Used to reduce sharp transitions in intensity
 - Reduce irrelevant detail in an image (e.g., noise)
 - Smooth the false contours that result from using an insufficient number of intensity levels in an image
- Filter kernels:
 - Box filter
 - Lowpass Gaussian filter
 - Order-statistic (nonlinear) filter

Box filter kernels

An array of 1's

$$M = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Normalizing constant




Box filters tend to favor blurring along perpendicular directions

Box filter example

Original image (1024x1024)

Kernel size: 3x3



Kernel size: 11x11




Kernel size: 21x21

Source: Fig. 3.33, Gonzalez

Gaussian filter kernels

When images with a high level of detail, with strong geometrical components

$$w(s, t) = G(s, t) = Ke^{-\frac{s^2 + t^2}{2\sigma^2}}$$



Gaussian filter example



Pattern image, 1024x1024



Gaussian filter size 21x21
 $\sigma = 3.5$



Gaussian filter size 43x43
 $\sigma = 7$

Source: Fig. 3.36. Gonzalez

Box vs Gaussian kernels



Original image



Box kernel, 21x21



Gaussian kernel, 21x21

Significantly less blurring

Note on Gaussian kernels

nothing to be gained by using a Gaussian kernel larger than $\lceil 6\sigma \rceil \times \lceil 6\sigma \rceil$

→ We get essentially the same result as if we had used an arbitrarily large Gaussian kernels





FIGURE 3.37 (a) Result of filtering Fig. 3.36(a) using a Gaussian kernels of size 43×43 , with $\sigma = 7$. (b) Result of using a kernel of 85×85 , with the same value of σ . (c) Difference image.

Shading correction using Gaussian filters




Order-statistic filters

- Nonlinear spatial filter
- Based on ordering (ranking) the pixels contained in the region encompassed by the filter
- Smoothing by replacing the value of the center pixel with the value determined by the ranking result
- Best-known filter:
 - Median filter
- Others:
 - Max filter
 - Min filter


Median filter

- Replaces the value of the center pixel by the median of the intensity values in the neighborhood of that pixel
- Excellent noise reduction:
 - Random noise
 - Impulse noise (salt-and-pepper noise)


Median filter in 1D



Median filter in 2D



Median filter example




X-ray image of a circuit board

Applied 19x19 Gaussian kernel
 $\sigma = 3$

Applied median kernel, 7x7

Source: Fig. 3.43, Gonzalez

Unsharp masking



Unsharp masking example



Original image, 600x259

Gaussian kernel, 31x31, $\sigma = 5$ 

Mask



Result of unsharp masking

Source: Fig. 3.49, Gonzalez

Spatial filtering with OpenCV

[Check the source code](#)